

OS

Introduction

Operating System

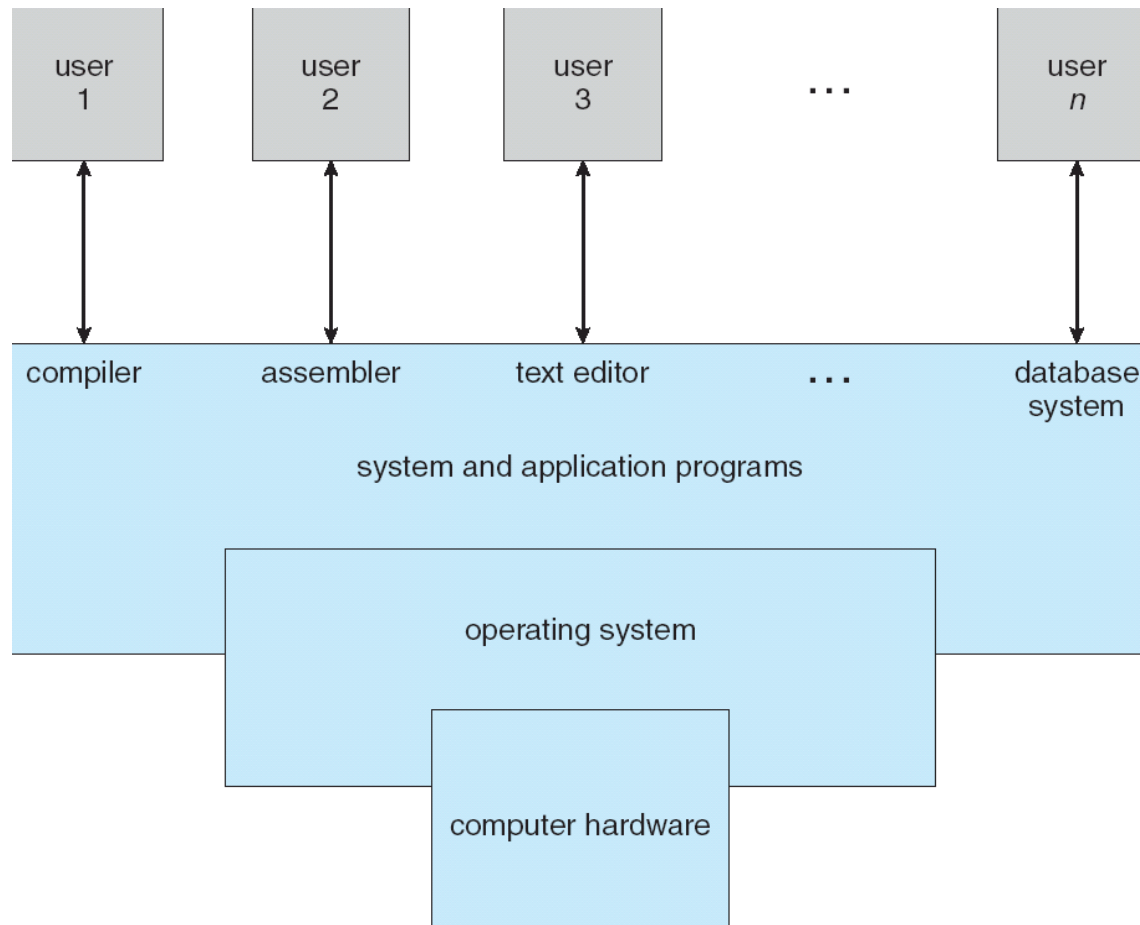
2

- A program that acts as an intermediary between a user and the computer hardware.
- Operating system goals:
 - ▣ Execute user programs and makes solving user problems easier.
 - ▣ Make the computer system convenient to use.
- Use the computer hardware in an efficient manner.

What operating systems do???

3

- Computer system can be divided into four components
 - ▣ Hardware – provides basic computing resources
 - CPU, memory, I/O devices
 - ▣ Operating system
 - Controls and coordinates use of hardware among various applications and users
 - ▣ Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - ▣ Users
 - People, machines, other computers



OS...

5

- OS is a **resource allocator**
 - ▣ Manages all resources
 - ▣ Decides between conflicting requests for efficient and fair resource use

- OS is a **control program**
 - ▣ Controls execution of programs to prevent errors

OS Definition...

6

.....No universally accepted definition

- “Everything a vendor ships when you order an operating system” is good approximation
- “The one program running at all times on the computer” is the **kernel**. Everything else is either a system program (ships with the operating system) or an application program.

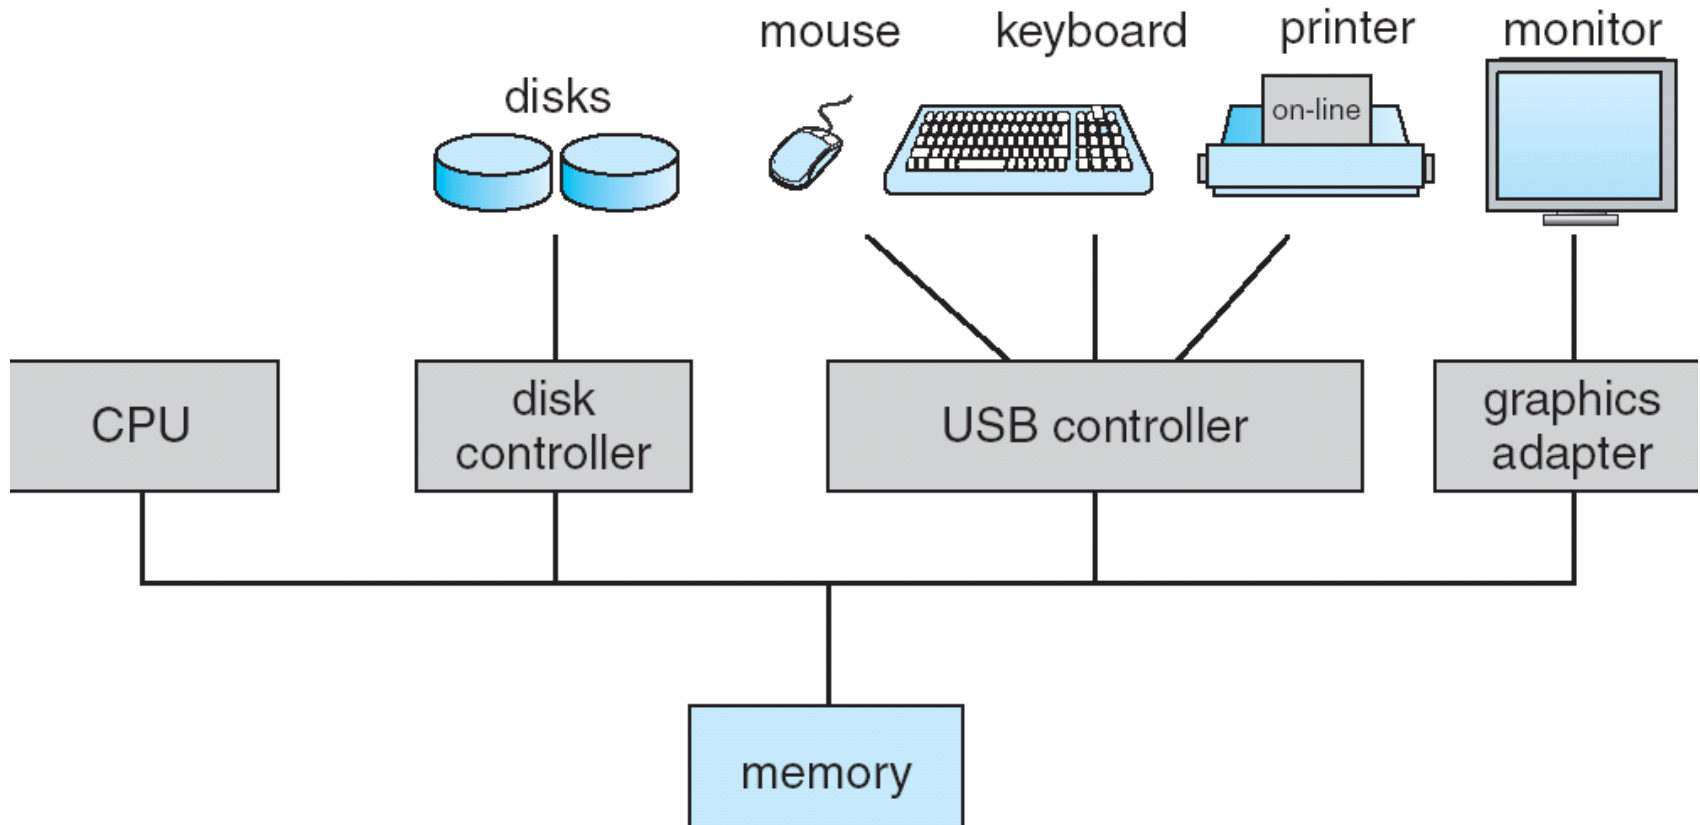
Computer Startup

7

- **bootstrap program** is loaded at power-up or reboot
 - ▣ Typically stored in ROM or EEPROM, generally known as **firmware**
 - ▣ Initializes all aspects of system
 - ▣ Loads operating system kernel and starts execution

Computer System Organization

8



Computer System Operation

- I/O devices and the CPU can execute concurrently.
- Each device controller is in charge of a particular device type.
- Each device controller has a local buffer.

- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller.
- Device controller informs CPU that it has finished its operation by causing an *interrupt*.

Functions of Interrupts

10

- Interrupt transfers control to the interrupt service routine generally, through the *interrupt vector*, which contains the addresses of all the service routines.
- Incoming interrupts are *disabled* while another interrupt is being processed to prevent a *lost interrupt*.
- An operating system is *interrupt* driven.

I/O Structure


11

- Device Driver for each Device Controller
- Controller starts the transfer of data from device to its buffer and interrupts Device Driver
- Device Driver returns control to the OS

Two I/O methods...

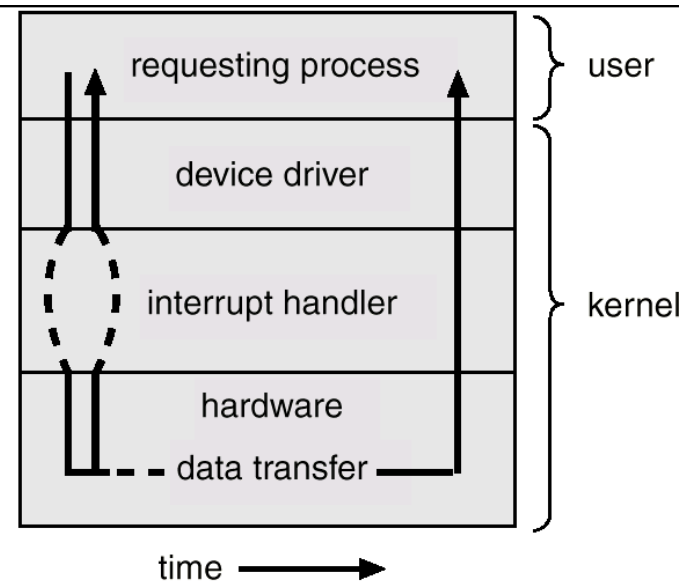
12

Asynchronous



(a)

Synchronous



(b)

Direct Memory Access Structure

13

- Used for high-speed I/O devices able to transmit information at close to memory speeds.
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention.
- Only one interrupt is generated per block, rather than the one interrupt per byte.

Storage Structure

14

- Main memory
 - ▣ only large storage media that the CPU can access directly.
- Secondary storage
 - ▣ extension of main memory that provides large nonvolatile storage capacity.

Storage Hierarchy


15

- Storage systems organized in hierarchy.
 - ▣ Speed
 - ▣ Cost
 - ▣ Volatility

- *Caching* – copying information into faster storage system; main memory can be viewed as a last *cache* for secondary storage.

Storage-Device Hierarchy


16



Caching

17

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - ▣ If it is, information used directly from the cache (fast)
 - ▣ If not, data copied to cache and used there.




Operating System Structure

18

- **Multiprogramming** needed for efficiency
 - ▣ Single user cannot keep CPU and I/O devices busy at all times
 - ▣ Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - ▣ One job selected and run via **job scheduling**
 - ▣ When it has to wait (for I/O for example), OS switches to another job
- **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - ▣ **Response time** should be < 1 second
 - ▣ Each user has at least one program executing in memory
⇒ **process**
 - ▣ If several jobs ready to run at the same time ⇒ **CPU scheduling**

Memory Layout for Multiprogrammed System

19



Operating-System Operations

20

- Interrupt driven by hardware
- Software error or request creates **exception** or **trap**
 - ▣ Division by zero, request for operating system service
- Other process problems include infinite loop, processes modifying each other or the operating system
- **Dual-mode** operation allows OS to protect itself and other system components
 - ▣ **User mode** and **kernel mode**
 - ▣ **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user

Process Management

21

- A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
 - ▣ CPU, memory, I/O, files
 - ▣ Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - ▣ Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - ▣ Concurrency by multiplexing the CPUs among the processes / threads

Process Management Activities

22

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling

Memory Management

23

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management determines what is in memory and when
 - ▣ Optimizing CPU utilization and computer response to users
- Memory management activities
 - ▣ Keeping track of which parts of memory are currently being used and by whom
 - ▣ Deciding which processes and data to move into and out of memory
 - ▣ Allocating and deallocating memory space as needed

OPERATING-SYSTEM STRUCTURES



OS Services

25

- **User interface** - Almost all operating systems have a user interface (UI)
 - ▣ Command-Line (CLI)
 - ▣ Graphics User Interface (GUI)
- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally
- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device.
- **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file Information, permission management

OS Services (cont...)

26

- **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
- **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

OS Services (cont...)

27

- **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
- **Accounting** - To keep track of which users use how much and what kinds of computer resources
- **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - ▣ **Protection** involves ensuring that all access to system resources is controlled
 - ▣ **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
 - ▣ If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.

User Operating System Interface - CLI

28

CLI allows **direct command entry**

- Sometimes implemented in kernel, sometimes by systems program
- Primarily fetches a command from user and executes it
 - Sometimes commands built-in, sometimes just names of programs

User Operating System Interface - GUI

29

- User-friendly **desktop** interface
 - ▣ Usually mouse, keyboard, and monitor
 - ▣ **Icons** represent files, programs, actions, etc
 - ▣ Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
- Many systems now include both CLI and GUI interfaces
 - ▣ Microsoft Windows is GUI with CLI “command” shell

System Calls

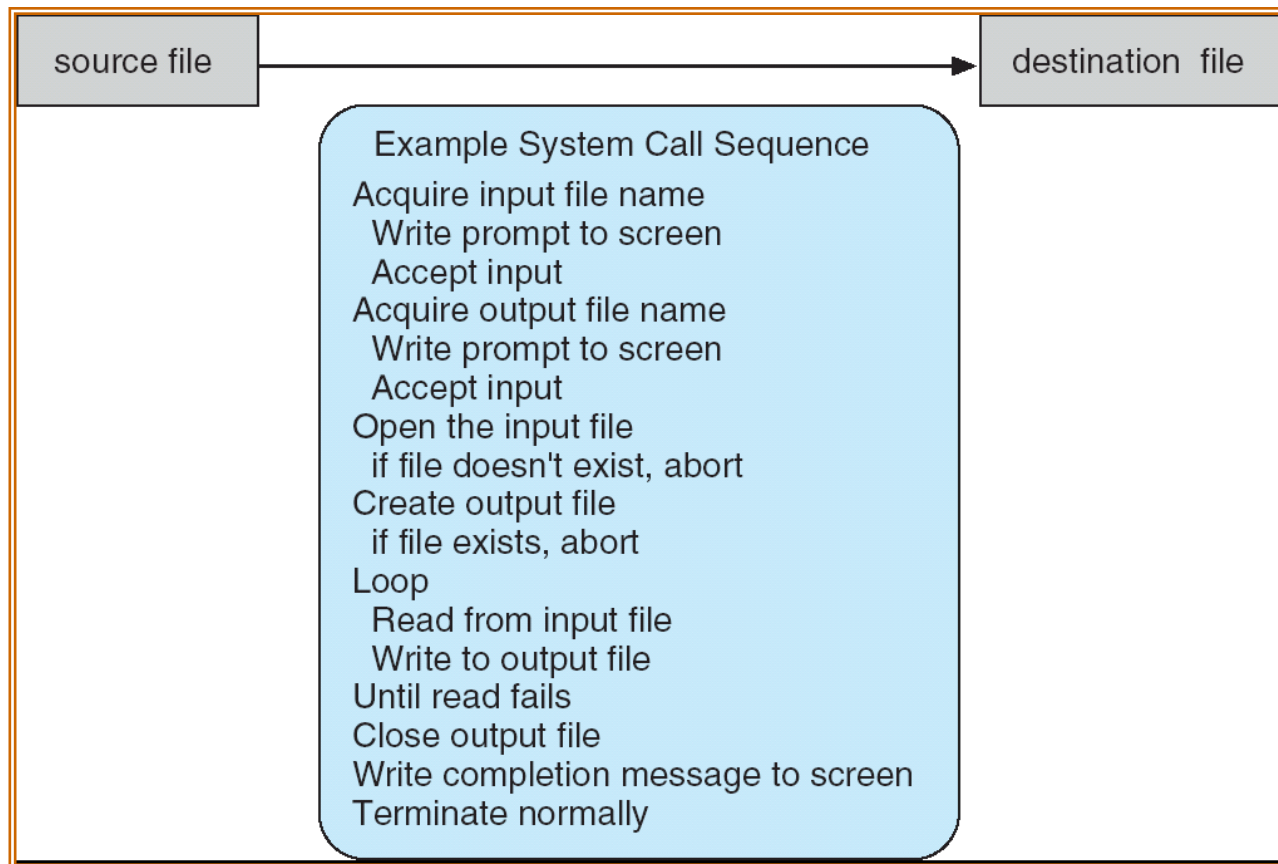
30

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level Application Program Interface (API) rather than direct system call use
- Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)

Example of System Calls

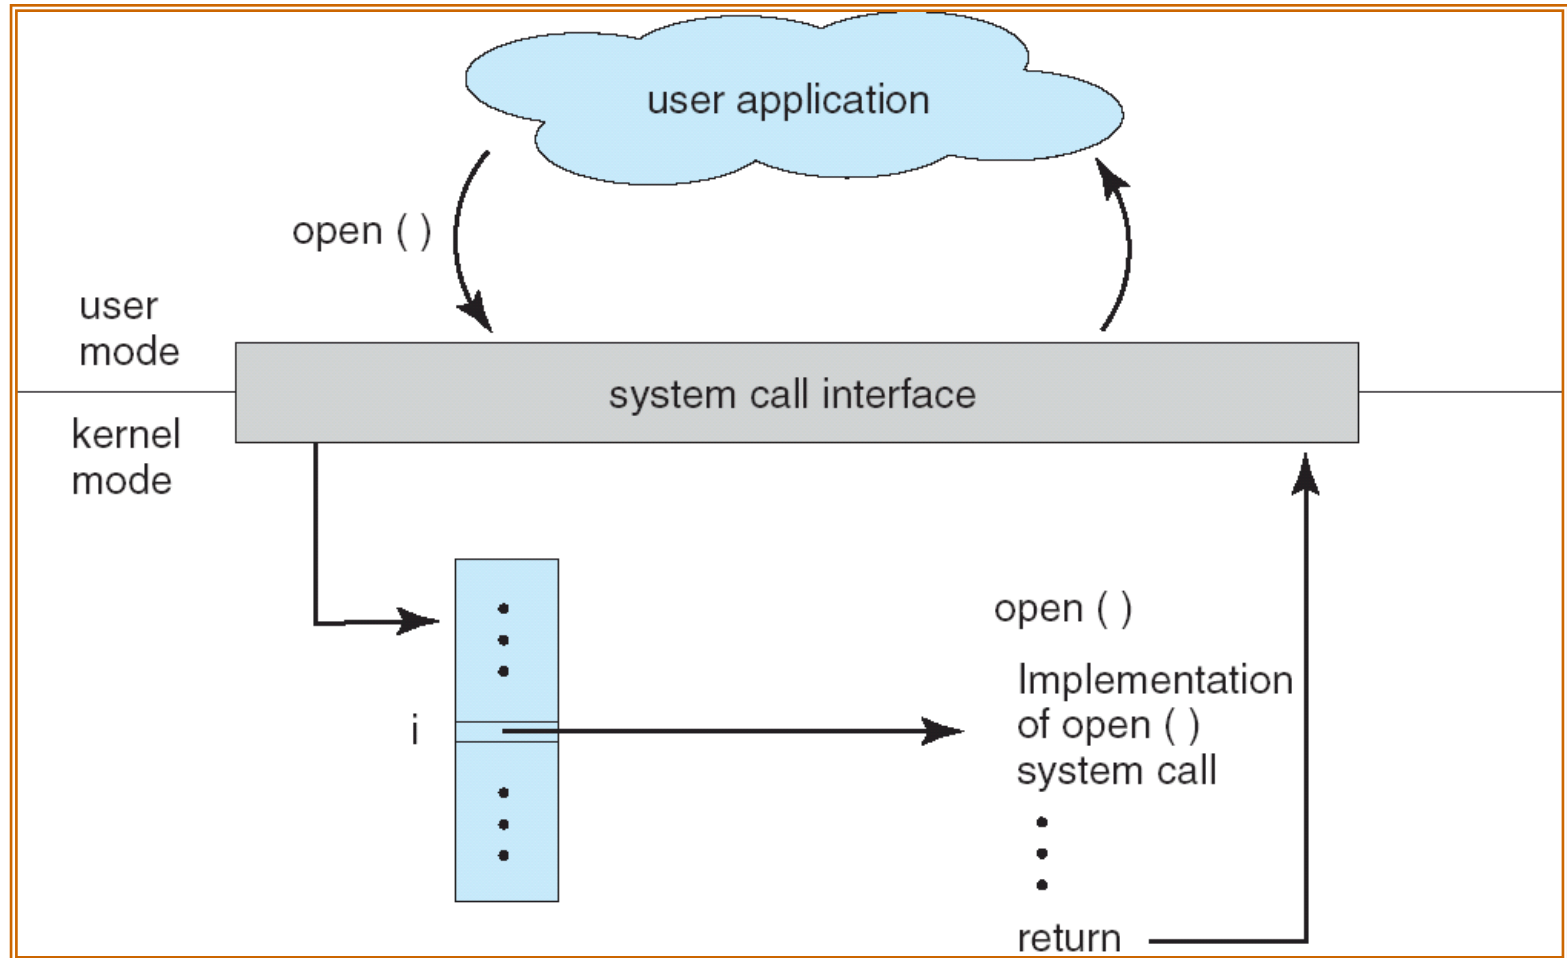
31

- System call sequence to copy the contents of one file to another file



API – System Call – OS Relationship

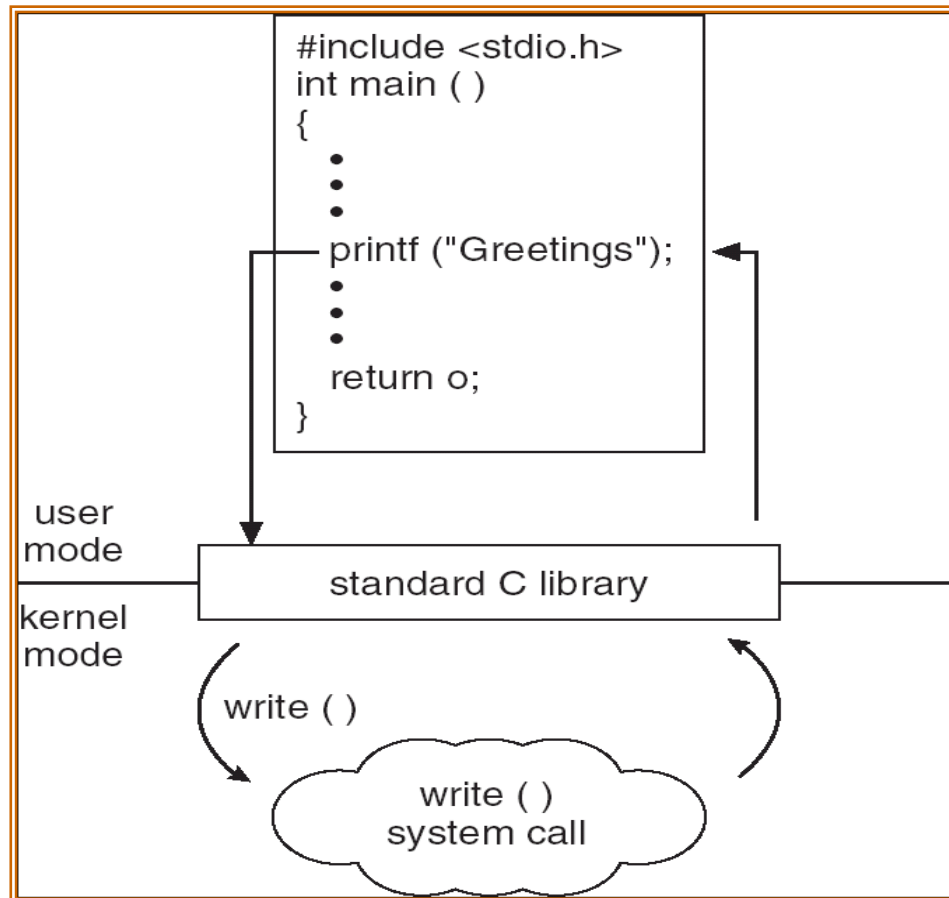
32



Standard C Library Example

33

- C program invoking printf() library call, which calls write() system call



Types of System Calls

34

- Process control
 - ▣ Create, terminate, wait, allocate, end, abort
- File management
 - ▣ Create, delete, open, close, read, write.
- Device management
 - ▣ Request, release, read, write
- Information maintenance
 - ▣ Get time, set time, get system data, set system data.
- Communications
 - ▣ Create ,delete communication, send receive messages

OS Design and Implementation

35

- Internal structure of different Operating Systems can vary widely
- Start by defining goals and specifications
- Affected by choice of hardware, type of system
- **User goals** and **System goals**
 - ▣ **User goals** – operating system should be convenient to use, easy to learn, reliable, safe, and fast
 - ▣ **System goals** – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient


Layered Approach

36

- The operating system **is divided into a number of layers** (levels), each built on top of lower layers.
 - ▣ The bottom layer (layer 0) is the hardware
 - ▣ The highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers

Layered Operating System

37



Operating System Generation

38

- Operating systems are designed to run on any of a class of machines; the system must be configured for each specific computer site
- **Booting** – starting a computer by loading the kernel
- **Bootstrap program** – code stored in ROM that is able to locate the kernel, load it into memory, and start its execution

PROCESSES




Process Concept

40

- An operating system executes a variety of programs concurrently.
- **Process – a program in execution**; process execution must progress in sequential fashion
- System – collection of processes: operating system processes executing system code & user processes executing user code.
- A process includes:
 - program counter
 - stack
 - data section

Process in Memory

41



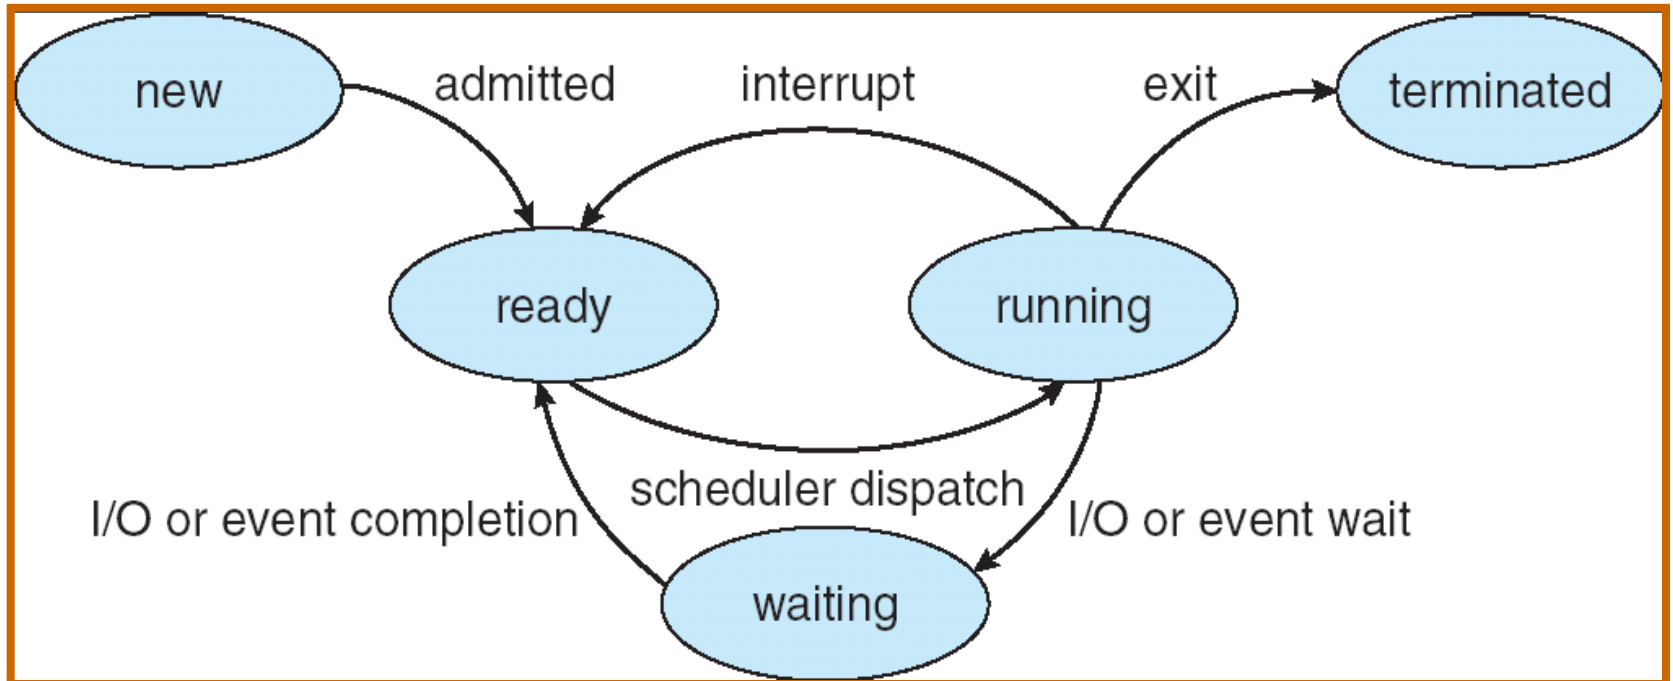
Process State

42

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a process
 - **terminated**: The process has finished execution

Process State Diagram

43



Process Control Block (PCB)

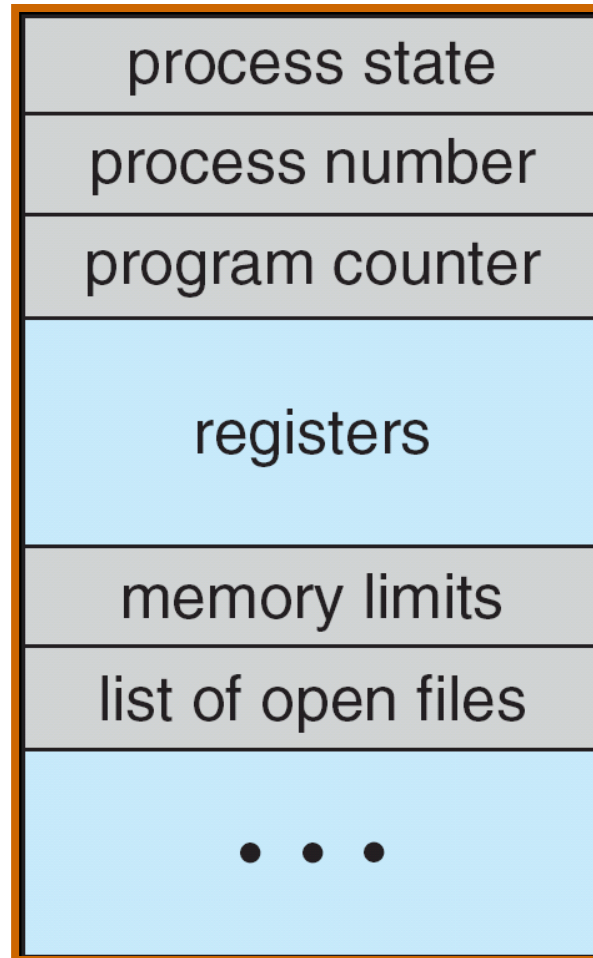
44

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

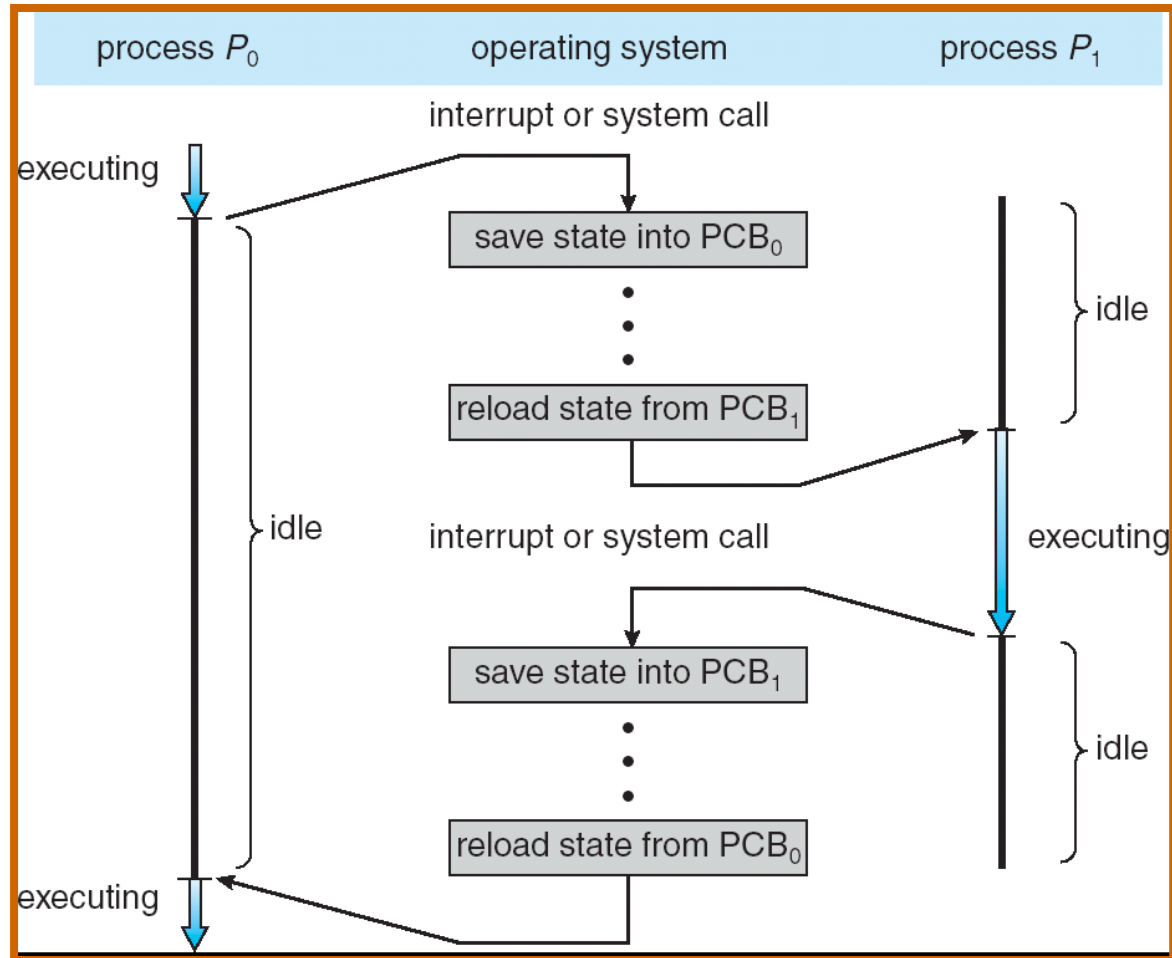
Process Control Block (PCB)

45



CPU Switch From Process to Process

46



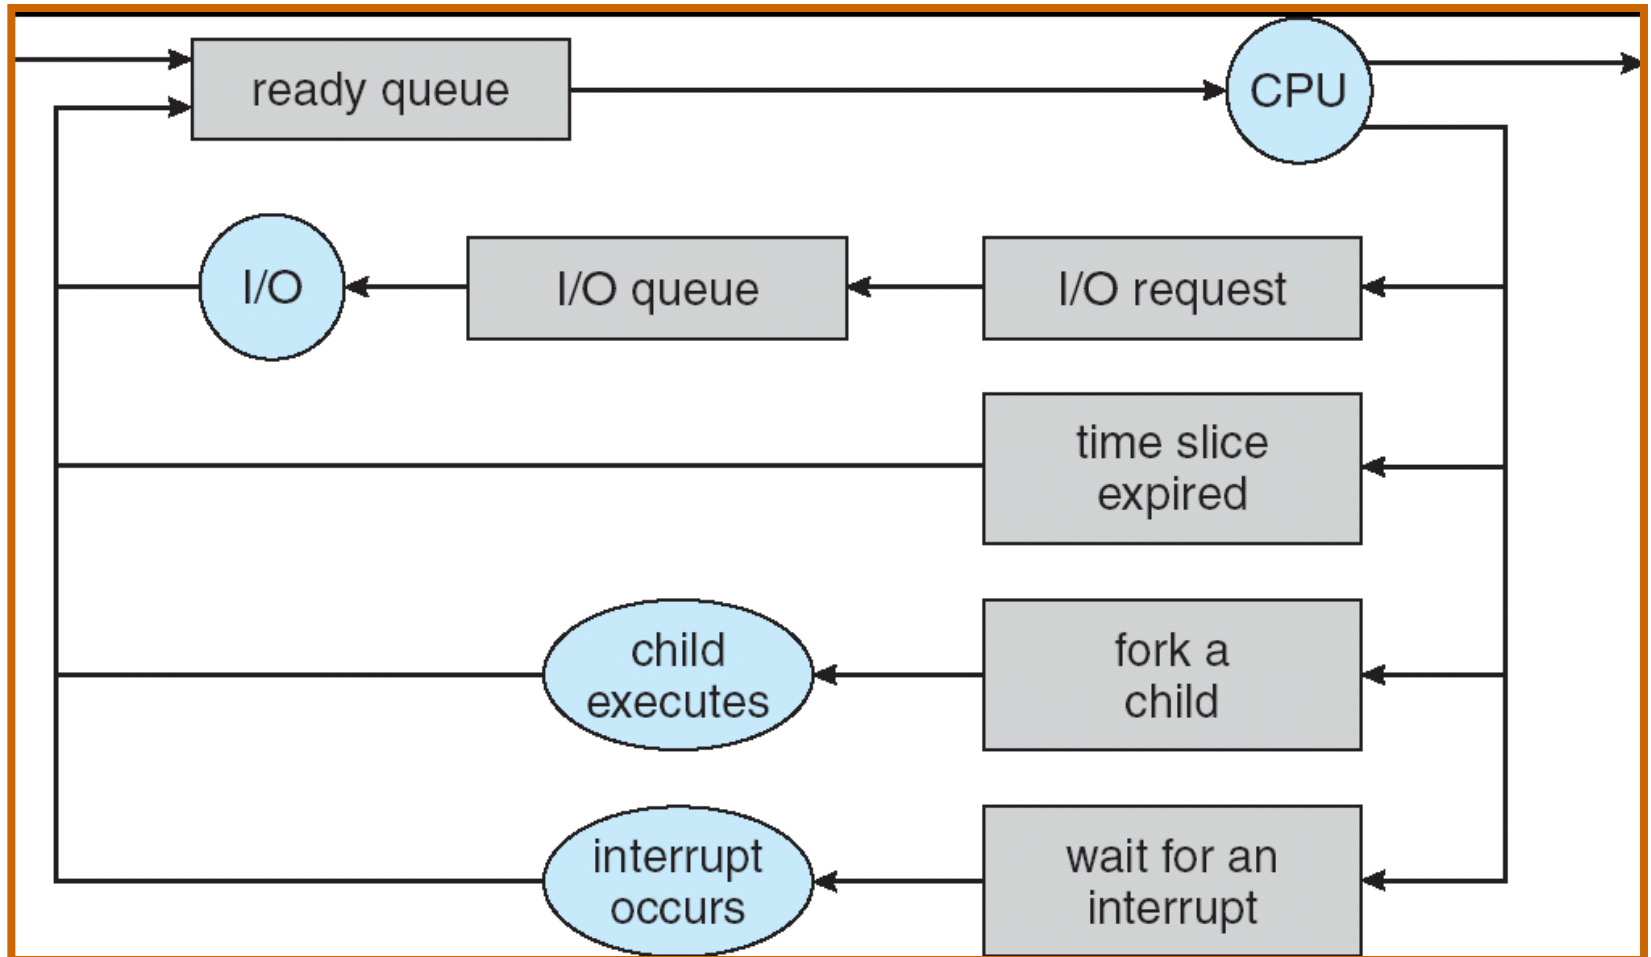
Process Scheduling Queues

47

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues

Representation of Process Scheduling

48



Schedulers

49

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked very infrequently (seconds, minutes) (may be slow)
 - The long-term scheduler controls the *degree of multiprogramming*
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
 - Short-term scheduler is invoked very frequently (milliseconds) ? (must be fast)

Schedulers (Cont.)

50

- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

Context Switch

51

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support

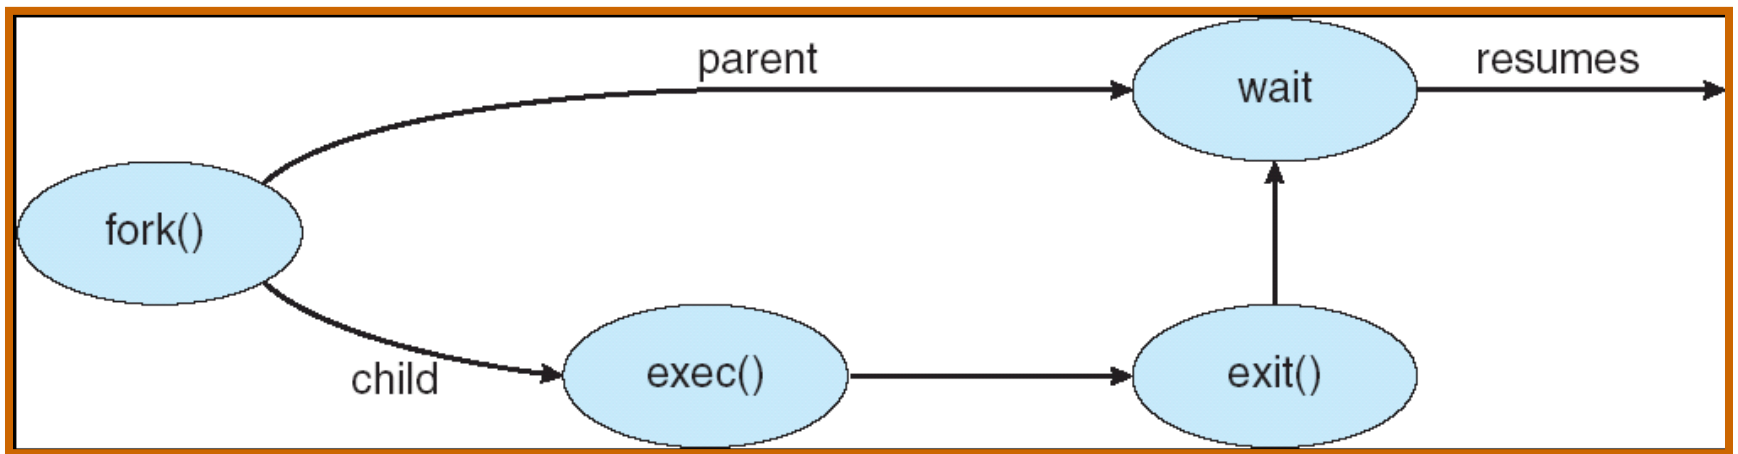
Process Creation

52

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation

53



Process Termination

54

- Process executes last statement and asks the operating system to delete it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - ▶ Some operating system do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*

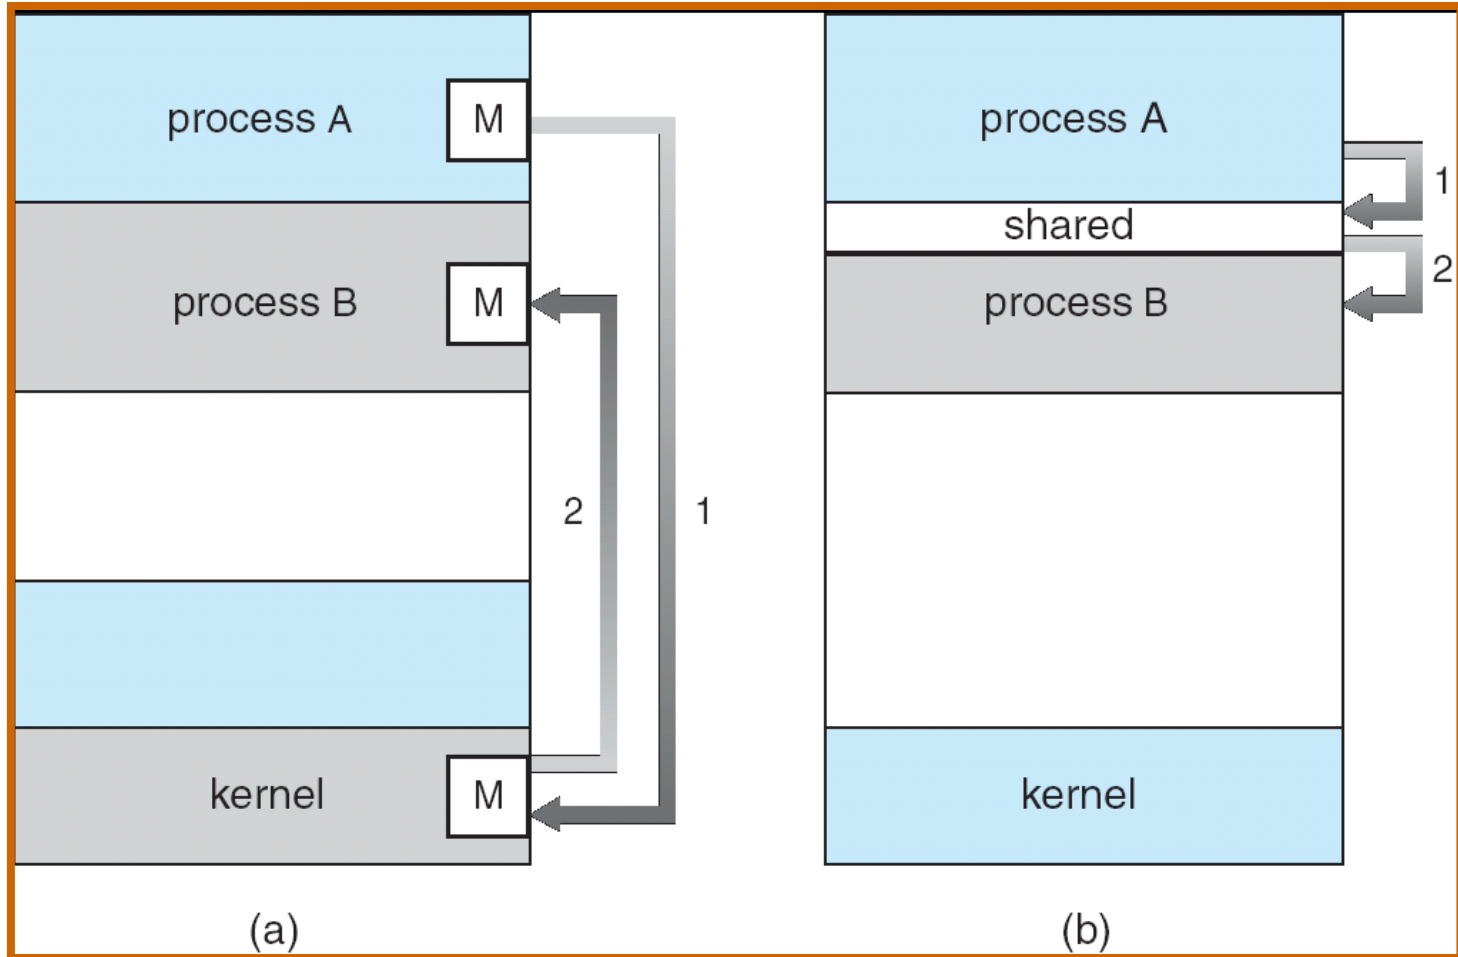
Interprocess Communication (IPC)

55

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Communications Models

56



Direct Communication

57

- Processes must name each other explicitly:
 - **send** (*P, message*) – send a message to process P
 - **receive**(*Q, message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

58

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

59

■ Operations

- create a new mailbox
- send and receive messages through mailbox
- destroy a mailbox

■ Primitives are defined as:

send(*A, message*) – send a message to mailbox A

receive(*A, message*) – receive a message from mailbox A

Indirect Communication

60

■ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 , sends; P_2 and P_3 receive
- Who gets the message?

■ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

61

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null

CPU SCHEDULING



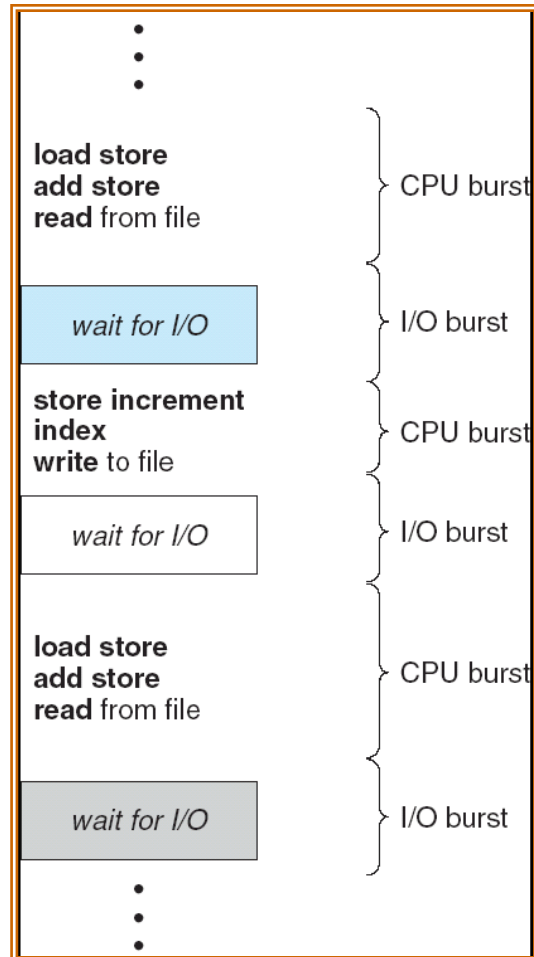
Why Scheduling???

63

- Maximum CPU utilization obtained with multiprogramming
- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait
- CPU burst distribution

Alternating Sequence of CPU And I/O Bursts

64



CPU Scheduler

65

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is *nonpreemptive*
- All other scheduling is *preemptive*

Dispatcher

66

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ▣ switching context
 - ▣ switching to user mode
 - ▣ jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

67

- CPU utilization
 - ▣ Keep CPU as busy as possible
- Throughput
 - ▣ Number of processes that are completed per time unit
- Turnaround time
 - ▣ Interval from time of submission to completion of a process
- Waiting time
 - ▣ Time spent in waiting queue
- Response time
 - ▣ Time from submission to first response to a request

Optimization Criteria

68

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

Scheduling algorithms

69


- First come first served
- Shortest job first
- Priority based scheduling
- Round robin

First-Come, First-Served (FCFS) Scheduling

70

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$


FCFS Scheduling (Cont.)

71

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- *Convoy effect* short process behind long process

Shortest-Job-First (SJF) Scheduling

72


- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - ▣ nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
 - ▣ preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes

Example of Non-Preemptive SJF

73

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)




- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

74

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

□ SJF (preemptive)



□ Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Priority Scheduling

75

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - ▣ Preemptive
 - ▣ nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process

Round Robin (RR)

76


- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.

Example of RR with Time Quantum = 20

77

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24


□ The Gantt chart is:



□ Typically, higher average turnaround than SJF, but better *response*

Time Quantum and Context Switch Time

78



Real-Time Scheduling

79

- ***Priority based preemptive scheduling***
- ***Hard real-time systems*** –
 - ▣ required to complete a critical task within a guaranteed amount of time
- ***Soft real-time computing*** –
 - ▣ requires that critical processes receive priority over less fortunate ones